

Business Objectives and Complex Portfolio Optimization

Peter Carl
Brian Peterson
Kris Boudt

Overview

- ▶ Touch on challenges in portfolio optimization
- ▶ Introduce the *PortfolioAnalytics* package
 - ▶ Demonstrate main functions and outputs
 - ▶ Describe implementation and usage issues
- ▶ Work through some examples - some “simple,” some more complex

Portfolio Optimization Distilled

- ▶ Markowitz (1952) described an investor's objectives as:
 - ▶ maximizing some measure of gain while
 - ▶ minimizing some measure of risk.
- ▶ Many approaches follow Markowitz and use mean return and standard deviation of returns for “risk”.
- ▶ Real investors often have more complex objectives.
- ▶ R contains a multitude of methods for solving optimization problems, most are not specific to finance.

Which Optimizer Should I Use?

- ▶ Linear, quadratic or conical objectives are better addressed through a package like RMetrics' *fPortfolio*.
 - ▶ *fPortfolio* will be faster for those problems.
 - ▶ See *Portfolio Optimization with R/Rmetrics*, by Diethelm Würtz, Yohan Chalabi, William Chen, Andrew Ellis.
- ▶ Many business objectives do not fall into those categories...
 - ▶ ...and brute force solutions are often intractable
 - ▶ Unconstrained, our example has over 68 million possible solutions

Frustrations with Optimization

- ▶ Users familiar with classic optimization describe a variety of problems:
 - Too many objectives
 - Wrong weighting of objectives
 - Too many parameters (too many assets)
 - Weights float to zero
 - Hard to understand what it's doing, or when it's broken
 - Too few solutions
 - Unrealistic expectations
 - Wrong optimization method for the job
 - “Worthless results”
- ▶ Most would prefer to be approximately correct rather than precisely wrong

You want to do *what*?

- ▶ Construct a portfolio that:
 - ▶ maximizes return,
 - ▶ with per-asset conditional constraints,
 - ▶ with a specific univariate risk limit,
 - ▶ while minimizing component risk concentration,
 - ▶ and limiting drawdowns to a threshold value.
- ▶ Not a quadratic (or linear, or conical) problem any more.

About *PortfolioAnalytics*

- ▶ *PortfolioAnalytics* focuses on providing numerical solutions for portfolios with complex constraints and objective sets comprised of any R function.
- ▶ Unifies the interface into different numeric optimizers, while preserving the flexibility to define any kind of objective and constraints.
- ▶ Provides a framework for managing different sets of portfolio constraints for comparison through time
 - ▶ Min risk, Equal risk, Equal weight, Position limits...
 - ▶ Supports regular and flexible rebalancing
- ▶ Builds intuition about optimization through visualization

About *PortfolioAnalytics*

- ▶ Currently implements a front-end to two analytical solvers, Differential Evolution and Random Portfolios
- ▶ Available on R-forge in the *ReturnAnalytics* project
 - ▶ `install.packages("PortfolioAnalytics", repos = "http://r-forge.r-project.org")`
- ▶ Work in progress, use $v \geq 0.5$, $rev \geq 1674$
- ▶ Functions are very compute intensive – even simple objectives may take a while (hours) to run on your netbook.
- ▶ Standard disclaimers apply: no warrantee, guarantees, etc.

About Random Portfolios

- ▶ At R/Finance 2009 and in multiple papers, Pat Burns describes using Random Portfolios to evaluate performance.
 - ▶ From a portfolio seed, generate random permutations that meet your constraints on the weights of each asset.
- ▶ Random Portfolio sampling can help provide insight into the goals and constraints of the optimization.
 - ▶ Aims to cover the 'edge case'(min/max) constraints almost completely, and evenly cover the 'interior' portfolios.
 - ▶ Useful for finding the search space for an optimizer.
 - ▶ Allows arbitrary number of samples.
 - ▶ Allows massively parallel execution.

About Differential Evolution

- ▶ Differential Evolution is a very powerful, elegant, population based stochastic function minimizer.
 - ▶ Continuous, evolutionary optimization.
 - ▶ Uses real-number parameters.
- ▶ Package *DEoptim* provides the algorithm in R.
 - ▶ Implementation of the algorithm distributed with the book:
 - ▶ *Differential Evolution - A Practical Approach to Global Optimization* by Price, K.V., Storn, R.M., Lampinen J.A, Springer-Verlag, 2005.
 - ▶ Thanks to R authors David Ardia and Katharine Mullen!

Load some packages, data


- ▶ Using returns for indexes representing four asset classes.
 - ▶ US Stocks, US Bonds, Int'l Stocks, Commodities.
 - ▶ 10 years of monthly total returns.

```
> library(PortfolioAnalytics)
> data(indexes)
> head(indexes)
```

	US Bonds	US Equities	Int'l Equities	Commodities	US Tbill	Inflation
2000-01-31	-0.0026	-0.0529	-0.0677	0.0674	0.0044	0.0058
2000-02-29	0.0120	-0.0193	0.0264	0.0588	0.0046	0.0091
2000-03-31	0.0135	0.0891	0.0375	-0.0117	0.0048	0.0000
2000-04-30	-0.0038	-0.0310	-0.0553	-0.0092	0.0049	0.0011
2000-05-31	-0.0002	-0.0209	-0.0248	0.1007	0.0048	0.0057
2000-06-30	0.0199	0.0241	0.0379	0.0665	0.0049	0.0023

Asset Returns and Risk

Returns



► Note:

- Huge discrepancy in risk and returns
- Obvious co-kurtosis and outlier effects
- Correlations increase markedly on negative shocks
- Generated with `charts.BarVaR`

Use an Equal Weight Benchmark

- ▶ Why an Equal Weight portfolio benchmark?
 - ▶ An equal weight portfolio provides a benchmark to evaluate the performance of an optimized portfolio against.
 - ▶ Each asset in the portfolio is purchased in the same quantity at the beginning of the period.
 - ▶ The portfolio is rebalanced back to equal weight at the beginning of each quarter.
 - ▶ Implies no information about return or risk.
- ▶ Helps answer questions a portfolio manager might ask:
 - ▶ Is the re-weighting adding or subtracting value?
 - ▶ Do we have a “useful” view of return and risk?

Calculate Equal Weight Benchmark

```
> dates=c(as.Date("1999-12-31"), time(indexes[ endpoints(indexes,
on="quarters") ]))
> weights = xts(matrix(rep(1/4,39*4), ncol=4), order.by=dates)
> colnames(weights)= colnames(indexes[,1:4])
> head(weights)
```


	US Bonds	US Stocks	Int'l Stocks	Commodities
1999-12-31	0.25	0.25	0.25	0.25
2000-03-31	0.25	0.25	0.25	0.25
2000-06-30	0.25	0.25	0.25	0.25
2000-09-30	0.25	0.25	0.25	0.25
...				

```
> EqWgt = Return.rebalancing(indexes[,1:4],weights)
> head(EqWgt)
```

	portfolio.returns
2000-01-31	-0.01395000
2000-02-29	0.02026640
2000-03-31	0.02976144
2000-04-30	-0.02482500
...	

Monthly returns for a quarterly-rebalanced portfolio

About the Equal Weight Portfolio



► This was a difficult period for this long-only portfolio.

► Annualized Return: 1.6%

► Annualized Std Dev: 12.3%

► Annualized Sharpe ($R_f = 3\%$) : -0.1

► Worst Drawdown: -47%

Example: Mean-CVaR Portfolio

- ▶ Although this is a “simple” case, the objectives are real:
 - ▶ Maximize the return per unit of risk taken.
 - ▶ Hold assets long-only, with positions limited by policy.
 - ▶ Remain fully allocated at all times.
 - ▶ Rebalance the portfolio quarterly.
 - ▶ Define risk as “downside risk” rather than just volatility.
 - ▶ Consider skewness and kurtosis.
- ▶ But even these “simple” portfolio objectives turn out to be complex to evaluate.
- ▶ This “base case” *could* be re-formulated in a conical solver.

Specify Constraints

- ▶ A 'constraints' object is simply a container that holds some key parameters we care about.

```
> aConstraintObj <- constraint(assets =  
+ colnames(indexes[,1:4]),  
+ min=0.05, # minimum position weight  
+ max=c(0.85,0.5,0.5,0.3), # maximum position weight  
+ min_sum=0.99, # minimum sum of weights approx. 1  
+ max_sum=1.01, # maximum sum must also be 1 + epsilon  
+ weight_seq = generatesequence()) # possible weights for  
  random or brute force portfolios
```

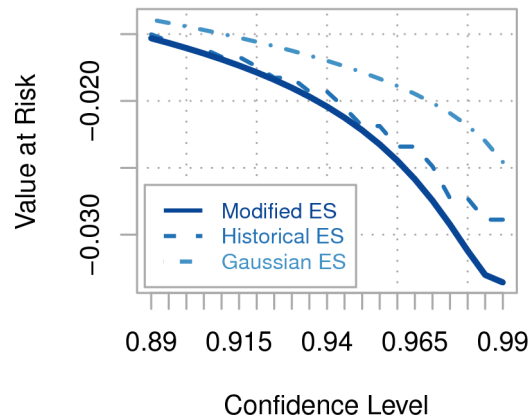
- ▶ Constraints may be specified for each asset in the portfolio individually.
 - ▶ Here we specify “box constraints” for min.

Specify Objectives

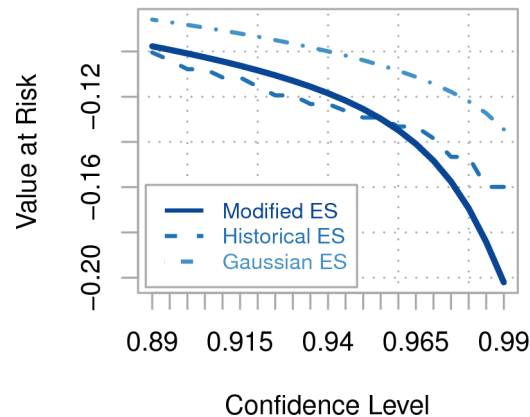
- ▶ Any function can be specified as an objective.
 - ▶ In this case, we use modified Conditional Value-at-Risk (CVaR or ES) as a univariate measure of portfolio risk.
 - ▶ Unlike Value-at-Risk (VaR), CVaR has all the properties a risk measure should have to be coherent and is a convex function of the portfolio weights.
 - ▶ We assume our return series is skewed and/or has excess kurtosis, so we use Cornish-Fisher estimates (or “modified”) of CVaR instead.
 - ▶ Usually convex, but can break down or be non-convex at extremes.
 - ▶ See ?CVaR in *PerformanceAnalytics*.

VaR Sensitivity

VaR Sensitivity for US Bonds



VaR Sensitivity for Int'l Equities




► 4-panel plot of VaR sensitivity from *PerformanceAnalytics*


► Modified CVaR (shown as ES) demonstrates a better fit for historical CVaR at lower confidence

► Breaks down at higher confidence levels

VaR Sensitivity for US Equities



VaR Sensitivity for Commodities



Specify Objectives

- ▶ Objectives are added to the constraint object. Here's a 'risk' objective:

```
> aConstraintObj <- add.objective(constraints =  
+ aConstraintObj, # our constraint object  
+ type="risk", # the kind of objective this is  
+ name="CVaR", # the function to minimize  
+ enabled=TRUE, # enable or disable the objective  
+ arguments=list(p=(1-1/12), clean='boudt')  
+ ) # parameters to pass to the CVaR function
```

- ▶ In this case, the CVaR function is portfolio-aware in that it takes returns and weights as arguments for evaluating each permutation.

Specify Objectives

- ▶ We need to pass the return series and the weighting vector for thousands of possible vectors, so we need to write a little wrapper to handle that:

```
> pamean <- function(n=12, R, weights, geometric=TRUE){  
+ # Trailing 12 month returns, annualized  
+ sum(Return.annualized(last(R,n),  
+ geometric=geometric)*weights)  
+ }
```

- ▶ Then we add the 'return' objective:

```
> aConstraintObj <- add.objective(constraints =  
+ aConstraintObj, type="return", name="pamean",  
+ enabled=TRUE, multiplier=-1,  
+ arguments = list(n=12))
```

Adding Portfolio Functions

- ▶ What we just did with “**pamean**” was define a new, arbitrary function for use by the optimization.
- ▶ Our example function is an portfolio annualized mean return function, but it could be any function you've written.
- ▶ If you name the return series “*R*” and the weights vector “*weights*”, the optimizer will populate these automatically.
- ▶ If your function has different arguments, you can specify them with the “*arguments*” parameter to **add.objective**.

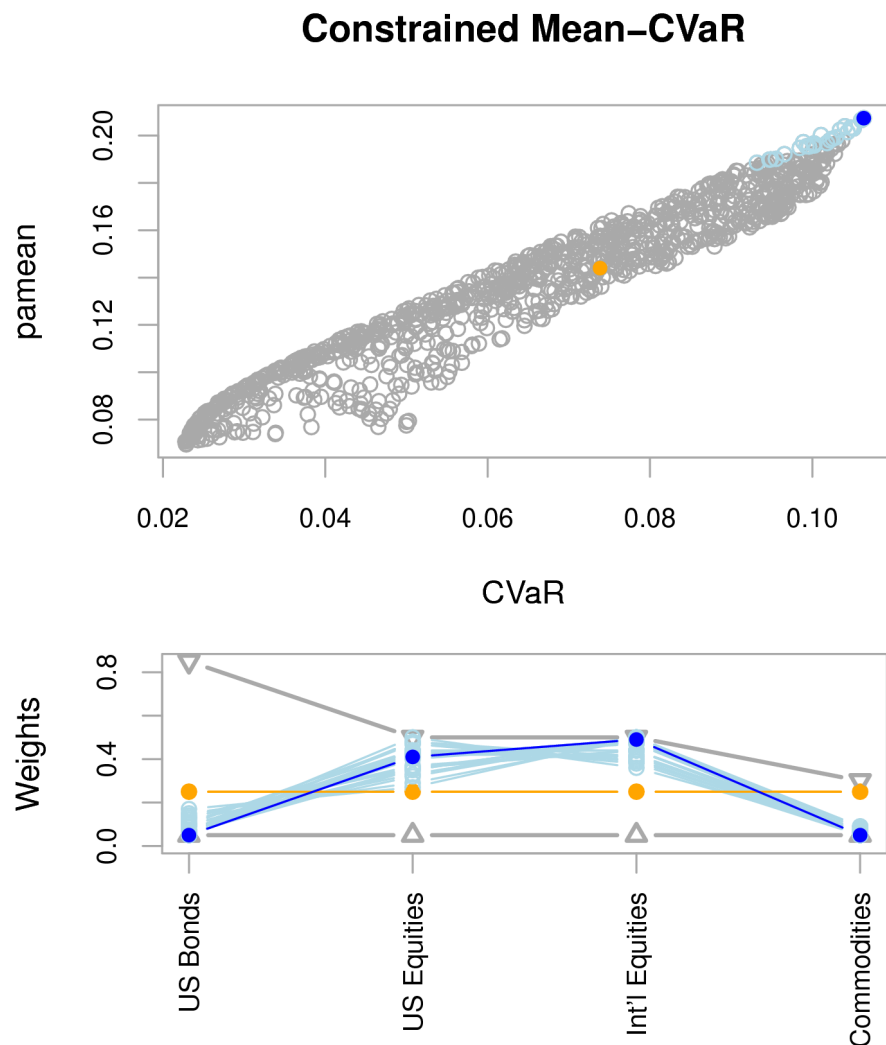
Specify Solver

- ▶ Generate sample portfolios for the most recent period:

```
> rndResult<-optimize.portfolio(R=indexes[,1:4],  
+ constraints=aConstraintObj, # our constraints  
+ optimize_method='random', # indicate solver to use  
+ search_size=1000, # number of portfolios to generate  
+ trace=TRUE, verbose=TRUE) # capture detail
```

- ▶ Our sample size should be about 1,000 portfolios.
 - ▶ For 4,050,000 possible combinations in the random portfolios with a step size of 1%, a 99% confidence and 2% error bands.
- ▶ Finds the optimal portfolio that maximizes the return per unit CVaR.

Mean-CVaR Results



- ▶ 1,000 unique, random portfolios within position constraints.
- ▶ Orange is the equal-weight portfolio.
- ▶ Blue is the optimal.
- ▶ Light blue shows 25 “near optimal” portfolios
- ▶ Weights are shown in the bottom panel.
- ▶ This is the default **plot** method for optimization

What Just Happened?

- ▶ The `optimize.portfolio` function manages the interface to the optimizer.
 - ▶ Instructs optimization backend to call `constrained_objective` for each target $w(eights)$.
- ▶ The `constrained_objective` function parses the constraints object and calls all the objective functions.
 - ▶ Applies penalty for failure to meet targets.
 - ▶ Summarizes the results in a single numerical output to be minimized.
- ▶ `optimize.portfolio.rebalancing` function manages the time loop and parallelization interface.

Mean-CVaR Through Time

- ▶ A few more parameters allow us to use the same constraint set through time:

```
> registerDoMC()  
  # get out more cores,  
  # this could be a different register* function  
  
> rndResults<-optimize.portfolio.rebalancing(R=indexes[,1:4],  
+ constraints=aConstraintObj, optimize_method='random',  
+ search_size=1000, trace=TRUE, # all the same as prior  
+ rebalance_on='quarters', # uses xts 'endpoints'  
+ trailing_periods=NULL, # calculates from inception  
+ training_period=36) # starts 3 years in
```

- ▶ Gives the optimal weights *each quarter* that maximize the return per unit CVaR (Minimum Risk).

Examine Results

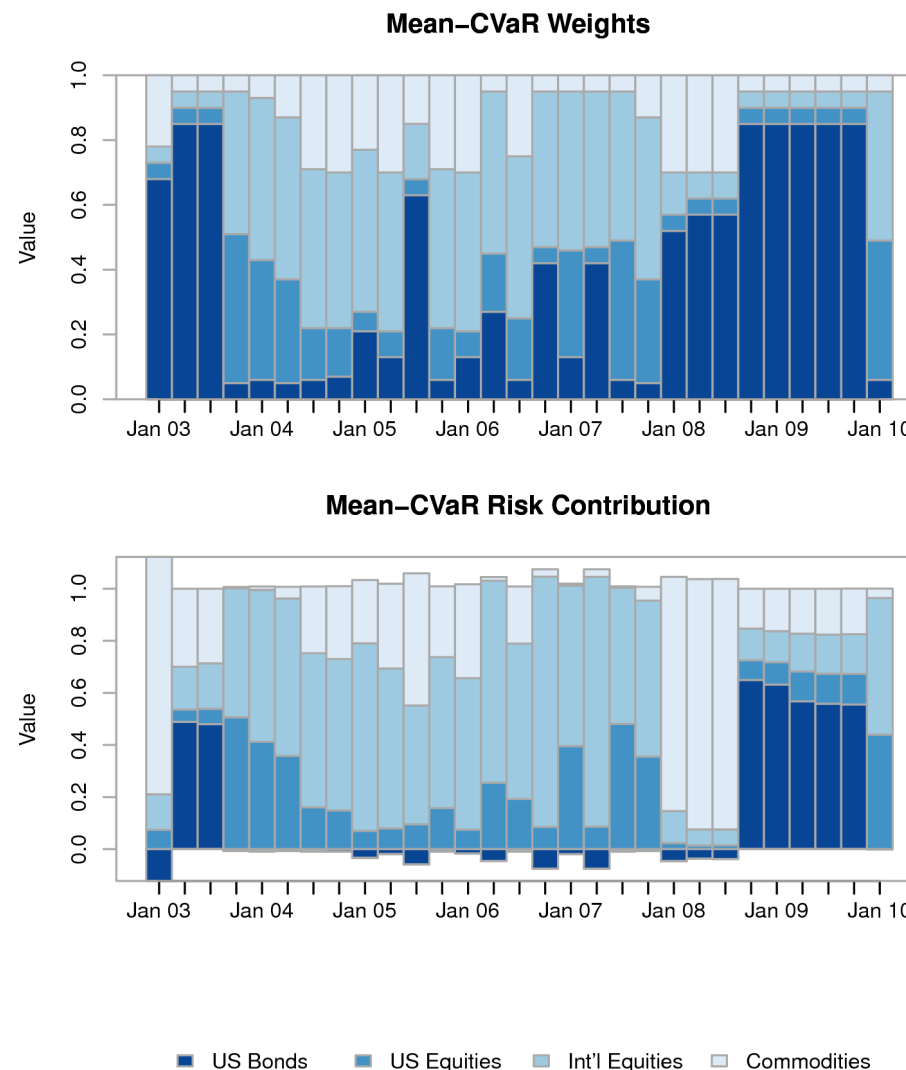
- ▶ Returns a list containing the optimal weights, some summary statistics, the function call, and optional trace information (turned off above).

```
> names(rndResults) # results organized by date
[1] "2003-12-31" "2004-03-31" "2004-06-30" "2004-09-30"
...
> names(rndResults[[1]]) # look at the first slot
[1] "weights"           "objective_measures" "call"
[4] "constraints"       "data_summary"      "elapsed_time"
[7] "end_t"
```


- ▶ **extractStats** function will pull out in-sample optimal portfolio statistics and weights for each rebalancing period
- ▶ Use **Return.portfolio** to calculate out of sample performance

Mean-CVaR Through Time

- ▶ Top panel shows weights through time.
- ▶ Second calculates contribution to portfolio CVaR through time.
- ▶ Components sum to 100%
- ▶ Diversifiers have contribution less than their portfolio weights, may have negative contributions



Mean-CVaR Through Time



- ▶ Controlling for risk improves performance...
- ▶ ... but lowers performance slightly during periods of stock out-performance

Example: Mean-CVaR w/ Risk Limit

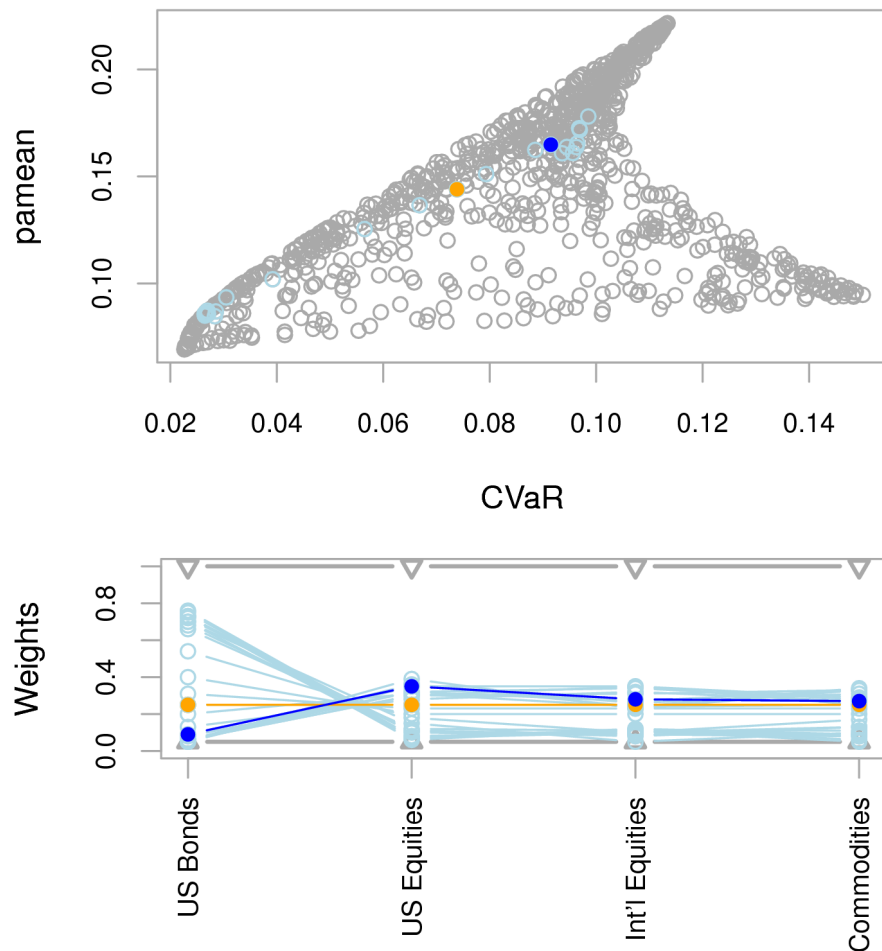
- ▶ Add another portfolio objective:
 - ▶ No asset can contribute more than 40% to the portfolio risk.
- ▶ We remove the original position limits, then add a risk budget objective with component risk limits:

```
> aConstraintObj$max <- rep(1,4)
> names(aConstraintObj$max) <- names(aConstraintObj$min)
> aConstraintObj <- add.objective(aConstraintObj,
+ type="risk_budget", name="CVaR", enabled=TRUE,
+ min_prisk=-Inf, # no negative limit
+ max_prisk=.4, # 40% contribution limit
+ arguments = list(clean='boudt', method="modified",
> p=(1-1/12))) # arguments for CVaR function

> rndResult2<-optimize.portfolio(R=indexes[,1:4],
+ constraints=aConstraintObj, optimize_method='random',
+ search_size=1000, trace=TRUE, verbose=TRUE)
+ ) # same as previous
```

Mean-CVaR w/ Risk Limit Portfolio

Mean-CVaR With Risk Limits

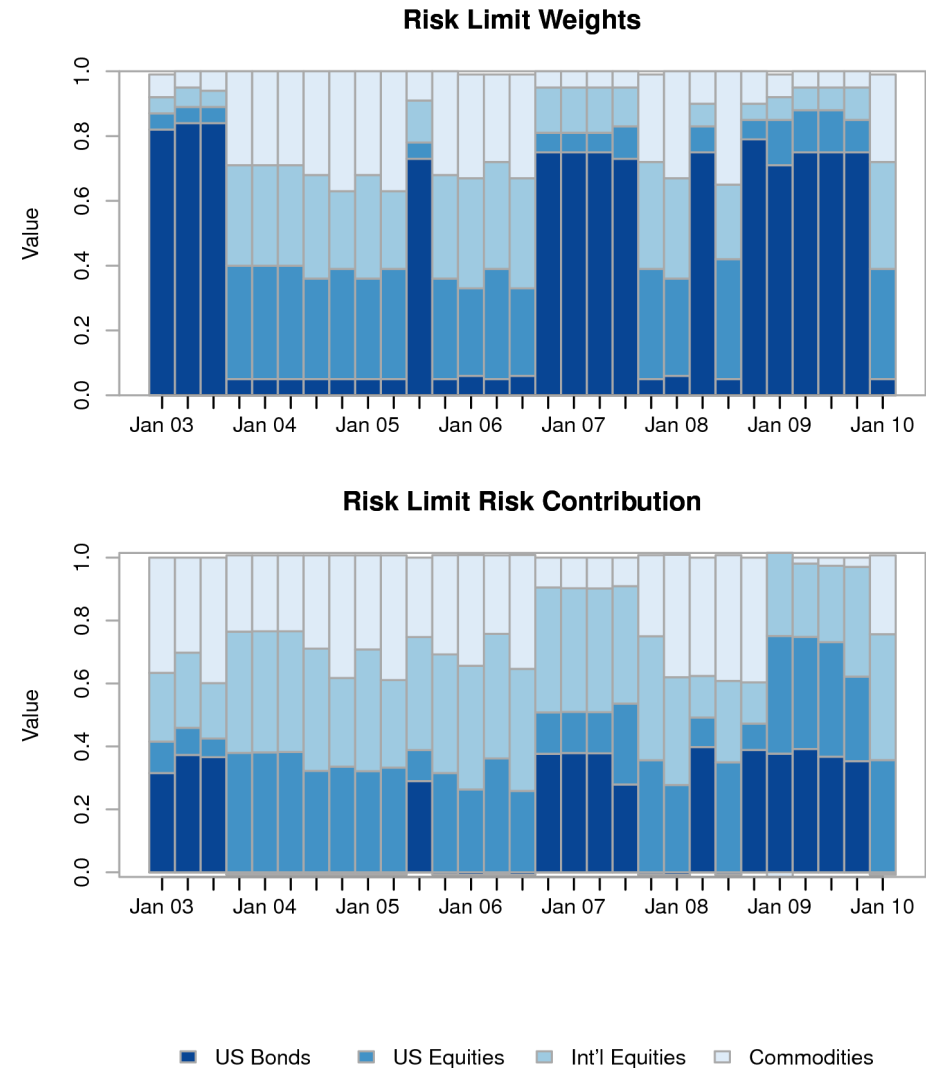


► Note:


- Difference in shape of the feasible space
- The optimal portfolio and its neighbors (near-optimal) are not on the outer hull
- Weights for the bond can vary over a wide range because their contribution to risk is so low

Mean-CVaR Risk Limit Portfolio

- ▶ Bond allocations emerge when equity and commodity risk increases
- ▶ Large bond allocations still contribute little in terms of portfolio risk



Mean-CVaR Risk Limit Performance



- ▶ These risk limits appear to constrain the portfolio more than the position limits did
- ▶ Better downside performance than Equal Risk, but worse than Mean-CVaR

Example: Equal Risk Portfolio

- ▶ Actually, this is the minimum component risk contribution concentration portfolio...
 - ▶ But it's easier to say “Equal Risk.”
- ▶ Why an Equal Risk portfolio?
 - ▶ Equal weight isn't necessarily balancing the portfolio risk
 - ▶ Equal risk looks to balance risk among the components of the portfolio.
 - ▶ Guard against estimation error on individual instruments (especially important on large real portfolios).
 - ▶ More likely to be “close” out of sample than traditional max/min objectives.

Specify the Equal Risk Constraints

- ▶ Build the constraint object:

```
> EqRiskConstr <- constraint(assets =  
+ colnames(indexes[,1:4]), min = 0.05,  
+ max = c(0.85,0.5,0.5,0.3), min_sum=1, max_sum=1,  
+ weight_seq = generatesequence())
```

- ▶ Add a “risk budget” objective and a “return” objective:

```
> EqRiskConstr <- add.objective(EqRiskConstr,  
+ type="risk_budget", name="CVaR", enabled=TRUE,  
+ min_concentration=TRUE, arguments = list(clean='boudt',  
+ p=(1-1/12)))  
> EqRiskConstr <- add.objective(constraints=EqRiskConstr,  
+ type="return", name="pamean", enabled=TRUE, multiplier=0,  
+ arguments = list(n=12))
```

- ▶ The zero multiplier in the return objective means that return will be calculated, but won't affect the optimization.

Use DEoptim

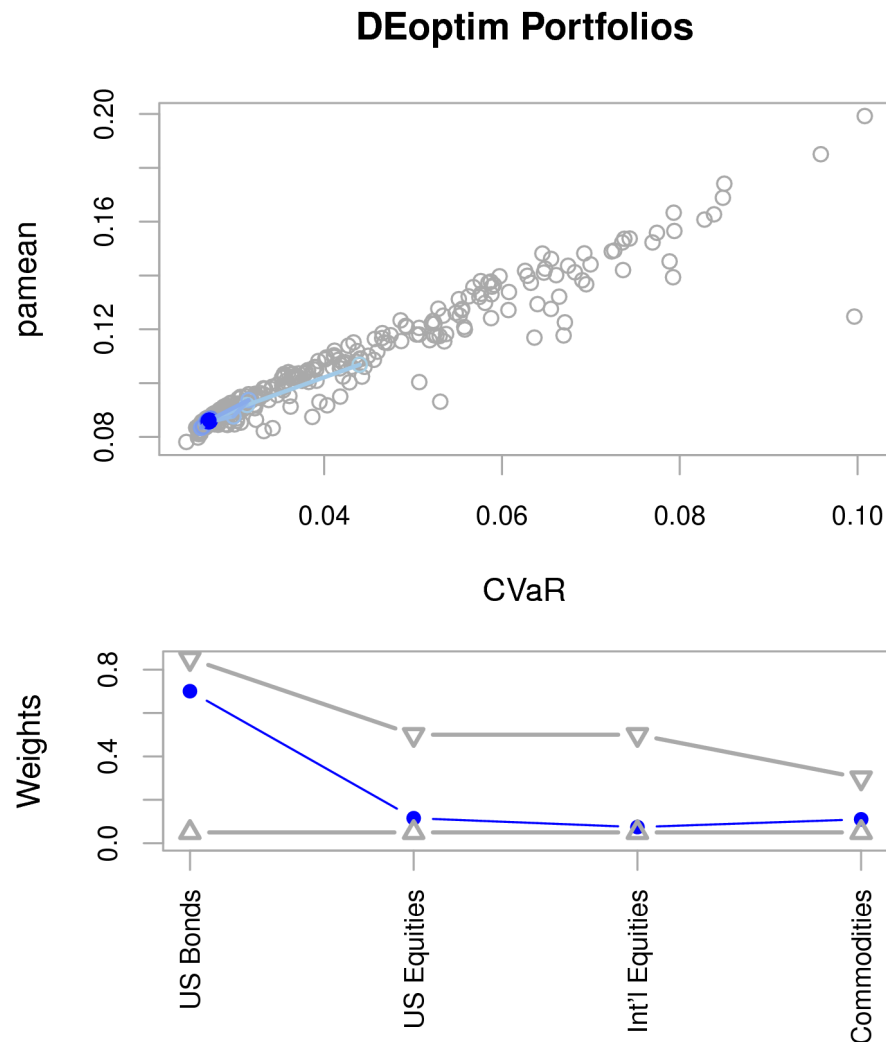
► Why DEoptim?

- All numerical optimizations are a tradeoff between speed and accuracy
- This space may well be non-convex in real portfolios
- DEoptim will get more directed with each generation, rather than the uniform search of random portfolios
- Allows more logical 'space' to be searched with the same number of trial portfolios for more complex objectives

► Specify DEoptim as the solver:

```
> EqRiskResultDE<-optimize.portfolio(R=indexes[,1:4],  
+ constraints=EqRiskConstr, optimize_method='DEoptim',  
+ search_size=2000, trace=TRUE, verbose=FALSE)
```

DEoptim Equal Risk Results



- ▶ DEoptim doesn't test many portfolios on the interior of the portfolio space
- ▶ Early generations search a wider space
- ▶ Later generations increasingly focus on the space that is near-optimal
- ▶ Random jumps are performed in every generation to avoid local minima


Run Optimizer Through Time

- Now we provide period information for rebalancing:

```
> EqRiskResultDEReбал <-  
+ optimize.portfolio.rebalancing(R=indexes[,1:4],  
+ constraints = aConstraintObj, # our constraints object  
+ optimize_method="DEoptim", # provide numeric sol'ns  
+ trace=FALSE, # set verbosity for tracking  
+ rebalance_on='quarters', # any xts 'endpoints'  
+ trailing_periods=NULL, # calculation from inception  
+ training_period=36, # starting period for calculation  
+ search_size=3000) # parameter to Deoptim, increase?
```


DEoptim Equal Risk Results

- ▶ Equal Risk objective provides a smoother return
- ▶ ... that underperforms the equal weight portfolio in most periods
 - ▶ given that it has no view on returns and suppresses weights in riskier assets
- ▶ ... but has a much smaller drawdown when things get ugly




Equal Risk Results

Equal Risk Weights



Equal Risk Risk Contribution



■ US Bonds ■ US Equities ■ Int'l Equities ■ Commodities


- ▶ We get a nearly-equal risk portfolio in almost all cases.
- ▶ 2Q2006 stands out as an exception
 - ▶ Maybe no feasible solution at the time?
 - ▶ Increase the search size?
 - ▶ Other DEoptim parameters?

Example: Mean-CDD

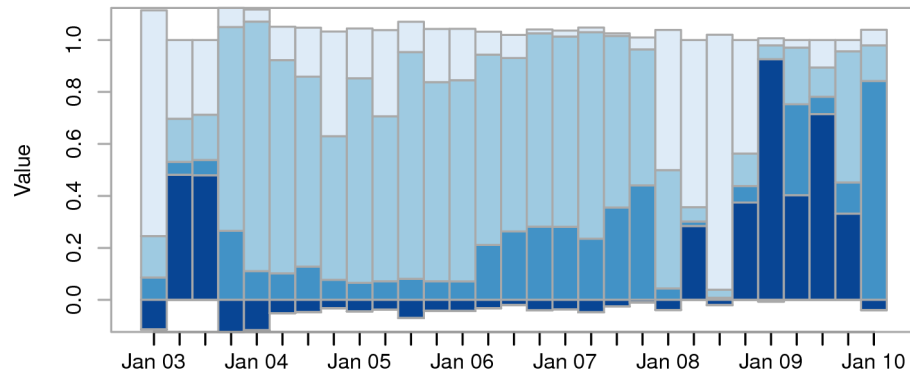
- ▶ Conditional Drawdown at Risk (CDD or CDaR) is the mean of the worst $p\%$ drawdowns proposed by Uryasev.
 - ▶ Another downside risk metric, but different than ES in that it does not assume independence of observations
 - ▶ Use the `name="CDD"` in `add.objective` to specify
- ▶ Qualitatively similar results to ES
 - ▶ Higher allocations to US Bonds and Int'l Stocks
- ▶ This fits in the broad class of “modified Sharpe” portfolio objectives

CDD with Return Objective

CDD w/ Return Obj. Weights




CDD w/ Ret. Obj. Risk Contribution



■ US Bonds ■ US Equities ■ Int'l Equities ■ Commodities

CDD Performance



Using more iron...

- ▶ *PortfolioAnalytics* uses Revolution's *foreach* package to run multiple optimizations to get a set of optimal portfolios
- ▶ DEoptim may only find a near-optimal solution, does it matter, or is it close enough?
- ▶ Examining the results of multiple runs – toward the central limit theorem
- ▶ `optimize.portfolio.parallel` will run an arbitrary number of portfolio sets in parallel
 - ▶ Develop confidence bands around your optimal solution
 - ▶ Show where the optimizer makes tradeoffs between assets
 - ▶ Highlight where you need larger number of portfolios or generations

Roadmap

- ▶ Additional portfolio analysis functions
- ▶ More portfolio-aware risk/return functions
- ▶ Bi-directional as.* functions for **portfolioSpec** in *fPortfolio*
- ▶ More testing, documentation, and demo code
- ▶ CRAN release

Contributions and Collaboration are Encouraged!

Getting Your Objectives Right

- ▶ What are *your* business objectives?
 - ▶ Most literature uses objectives too simple to be realistic
 - ▶ Most software reinforces this
- ▶ Random Portfolios help you see the shape of the feasible space
 - ▶ The scatter chart shows the space covered by the portfolios that meet your constraints
- ▶ Rebalancing periodically and examining out of sample performance will help refine objectives
- ▶ DEoptim and parallel runs are valuable as things get more complex